

Brief Technical Write-up

Tech Stack Choices

PDF Parsing : PyPDF2 - Lightweight, easy to integrate.

Text Chunking : LangChain RecursiveCharacterTextSplitter - Supports context overlap.

Embedding Model : TF-IDF (Scikit-learn) - Fast, local, and API-free.

Vector Search : FAISS - Efficient nearest-neighbor search for retrieval.

LLM (Cloud) : OpenAI GPT-3.5 - High-quality structured response generation.

LLM (Offline) : Ollama with LLaMA 3 - Fully offline generation, privacy-first.

Frontend UI : Streamlit - Simple interactive web interface.

Response Structuring Approach

- Prompt engineering ensures the model only responds using retrieved document context.
- Structured formatting (bullet points, tables, paragraphs) improves clarity.
- Fallback prompt added: "This information is not available in the document." to prevent hallucinations.
- Top-K (default 3) most relevant chunks are passed to the LLM for generation.

Challenges Faced & Solutions Implemented

Challenge: sentence-transformers errors

Solution : Replaced with TF-IDF + FAISS for embedding and retrieval.

Challenge: OpenAI v1+ API compatibility

Solution : Updated to openai.ChatCompletion (v1.0+ syntax).

Challenge: Local-only functionality

Solution : Integrated Ollama to run LLaMA 3 locally via HTTP POST requests.

Challenge: PDF format inconsistency

Brief Technical Write-up

Solution : Used PyPDF2 with fallback logic for non-standard formats.

Challenge: Streamlit path errors on Windows

Solution : Used `sys.path.append` and restructured folders for relative imports.